

Here's the fuller version, because the thing I'm doing with Suno makes a lot more sense once you know that there is an actual compositional system underneath the apparent nonsense.

I built a little experimental prompt app I call MY LIL GUYS FOR SUNO.

The name is intentionally stupid. The system is not.

The basic idea is that I am not primarily asking AI to "make me a song that sounds like genre A mixed with genre B mixed with genre C." I am much more interested in asking:

What happens if different parts of a composition are governed by different rules, different logics, different perceptual assumptions, and different kinds of cognition — and then all of those systems are forced to coexist inside one piece of music?

The central principle is:

DO NOT JUST BLEND THE INGREDIENTS.

GIVE THEM SEPARATE JURISDICTIONS AND FORCE THEM TO NEGOTIATE.

That distinction is extremely important.

If I say "combine barbershop harmony, Balkan rhythm, surf guitar, operatic vocals, industrial percussion, and doo-wop," a normal generative system tends to puree those things into an average. You get genre soup. Everything starts compromising with everything else until the weirdness becomes mostly cosmetic.

My system tries to prevent that.

Instead, I assign separate musical dimensions their own governing logic.

One system may control HARMONY.

One may control MELODY.

One may control RHYTHM.

One may control TIMBRE / ATMOSPHERE.

One may control PERFORMANCE ATTITUDE.

One may control VOCAL BEHAVIOR.

And usually there is some kind of ANCHOR or INVARIANT that remains recognizable while the rest of the song deforms around it.

So harmony might be operating according to one musical tradition while rhythm is operating according to a completely different one. Melody may obey a third logic. Vocal phrasing may follow something else entirely.

They are not supposed to merge politely.

They are supposed to retain enough autonomy that their interaction becomes structurally meaningful.

The harmony guy does not get to tell the rhythm guy how to do his job.

The rhythm guy does not get to domesticate the melody guy.

The vocal guy may decide that syllables are now percussion.

The timbre guy may suddenly turn the whole surface into cheap spring reverb and broken circuitry.

Meanwhile the anchor has one job: survive.

That is why I call them “the Lil Guys.”

They are not really personas in the normal sense. They are more like little cognitive jurisdictions.

Each one gets a job.

Each one gets rules.

Each one gets limits.

Each one gets some degree of independence.

And then I force them to share one song.

The interesting part is what happens at the borders.

A lot of the compositional energy comes from negotiation between systems that do not naturally agree.

For example:

The rhythmic system may insist on a repeating 7/8 cycle.

The melodic system may insist on long, floating phrases that do not care where the bar line is.

The harmony system may refuse conventional resolution.

The singer may be doing clipped rapid-fire technical narration.

The percussion may be physically simple and stupid — claps, stomps, dry hits — while the harmonic structure is mathematically obnoxious.

Now those systems have to coexist.

That tension is productive.

Instead of asking the model to “be weird,” I try to make weirdness emerge because the laws of the composition create conditions that are difficult to reconcile.

That is a huge part of this project:

THE WEIRDNESS SHOULD BE OPERATIONAL.

Not decorative.

I try very hard not to use words like “experimental,” “chaotic,” or “surreal” as if those words alone tell the system anything useful.

Instead I use operators.

Operators are actual transformations.

Things like:

hold A fixed while B migrates

compress one layer into a tiny duration while another stretches across a huge duration

make one system behave like another system without actually becoming it

preserve the anchor while destabilizing everything around it

force two incompatible temporal scales to coexist

alternate extreme precision with controlled collapse

remove an entire jurisdiction and force the remaining systems to continue without replacing it

return to the same anchor repeatedly, but each time in a more mutated state

make the vocal rhythm inherit the percussion pattern

make the percussion start behaving like speech

make the harmony freeze while every other dimension accelerates

make one performer’s output become the trigger for another performer’s entrance

That is much more interesting to me than saying “make it crazier.”

Because now the craziness has a cause.

It also means that when something interesting happens, I can sometimes identify WHY it happened.

That matters a lot.

I am trying to learn what kinds of generative mechanics actually produce interesting music instead of just collecting random outputs.

The app has also expanded way beyond just “assign these musical jobs.”

I’ve been building a library of different cognitive procedures — basically different ways of THINKING before the actual music prompt is generated.

These are some of the weirder parts of the system.

One cognitive procedure might take a property that normally belongs to one thing and reassign it to something that normally would not own it.

Not as metaphor.

Literally, inside the temporary logic of the prompt.

For example, instead of the singer having emotion, perhaps the meter has emotion.

Then you have to ask:

If meter is the thing that experiences emotional pressure, what changes?

Maybe rhythm becomes unstable specifically when emotional intensity increases.

Maybe tempo is not expressive, but meter itself fractures.

Maybe the vocalist remains emotionally flat while the bar structure panics.

That is one kind of Lil Guy thinking.

Another mechanism is AXIOM FRACTURE.

That means finding a hidden assumption that normally makes the obvious solution possible, negating that assumption, and then forcing the composition to succeed anyway.

So maybe the ordinary assumption is:

A melody needs continuity in order to remain recognizable.

Then I remove that assumption.

Now the system has to preserve melodic identity without continuity.

Maybe the melody survives only through recurrence of interval relationships across different instruments.

Maybe no phrase ever completes.

Maybe the listener reconstructs the melody from scattered pieces.

The point is not merely to make things difficult.

The point is to force the system to find a route that ordinary composition would not naturally select.

Another Lil Guy uses ATTENTIONAL SHADOW HARVESTING.

Instead of focusing on the obvious center of the music, it asks:

What is being ignored in order for the obvious musical idea to seem coherent?

What background condition is doing hidden work?

Maybe silence.

Maybe decay.

Maybe maintenance.

Maybe breathing.

Maybe bleed between instruments.

Maybe the count-in.

Maybe the recovery period after a phrase.

Then that neglected thing gets promoted into a governing condition.

Another cognitive mechanism is DESTRUCTIVE COMPRESSION.

This one strips a musical problem down until almost all recognizable genre identity is gone.

You remove instrumentation, cultural labels, stylistic decoration, era, aesthetic associations — anything that is not causally essential.

What remains is something like:

three systems compete for limited temporal space

one signal has to remain recognizable

interruptions increase over time

resource allocation shifts when density rises

That is the minimum machine.

Then the system temporarily solves that machine inside some completely unrelated domain.

Not because the unrelated domain looks cool.

Because its native rules provide a different structural solution.

Then the decorative imagery gets stripped away again.

Only the structural solution comes back into the music.

That helps prevent the AI from doing the lazy “this reminds me of X, therefore add some X-themed vocabulary” move.

Another mechanism is ERROR AXIOMATIZATION.

This one asks:

What is the nearest coherent WRONG way to classify this musical situation?

Not random error.

A structured near-miss.

Then:

What kind of world would have to exist for that wrong classification to actually be correct?

And once that temporary world is established, the system reasons rigorously inside it.

So maybe a chord progression is misclassified as a respiratory cycle.

Okay.

Now what would have to be true for harmony to literally behave like respiration?

Not “sound breathy.”

Actually behave according to intake, retention, exchange, release, oxygen debt, recovery.

Now harmonic pacing may follow pressure relationships rather than normal harmonic function.

That produces something much stranger than “add airy synths.”

There is another one called PRIMITIVE DELETION.

That one removes a foundational concept entirely.

Not reverses it.

Not weakens it.

Deletes it.

For example, maybe the system temporarily has no concept of “verse.”

Not “avoid verses.”

Verse as a conceptual primitive simply does not exist.

Now the song has to organize itself without quietly reinventing a verse under another name.

That absence becomes causally important.

Another one reasons backward from a fixed terminal state.

Instead of asking:

What could this song eventually become?

It decides what the final state definitely is and then reasons backward to determine what earlier conditions would have been necessary to make that ending inevitable.

Another treats apparently stable musical objects as processes rather than things.

A groove is not a groove.

It is a continuously maintained equilibrium produced by timing, repetition, muscle memory, expectation, energy input, and error correction.

Then instead of “adding another element,” the system modifies rates, gradients, friction, feedback, permeability, interruption frequency, or coupling strength.

That kind of thinking is useful because it changes what counts as an intervention.

I also have one that blocks direct representation of a target concept.

If I want something to express panic, the system may be forbidden from thinking directly in terms of panic.

It has to triangulate panic through unrelated axes:

temporal compression
resource allocation
attention fragmentation
information arrival rate
motor pressure
failure recovery

Then it generates from the intersection of those constraints instead of from stock “panic music” associations.

That helps avoid cliché.

Another one treats the composition as a feedback circuit.

Instead of saying “the percussion is causing chaos,” it asks:

What difference does the percussion create?

What system responds to that difference?

What new difference does THAT response create?

How does that eventually come back and alter the original percussion?

So you get actual recursive escalation instead of just “more layers.”

Another one replaces a deep conceptual primitive and recompiles all dependent concepts underneath it.

Another lets a foreign rule system infect the host composition.

Again, not stylistically.

Mechanically.

The foreign system might be ecology, bureaucracy, immune response, market behavior, swarm logic, fluid dynamics, or something else.

But I do not want the song to merely SOUND like that domain.

I want the operational lifecycle of that domain to colonize the music.

If it is an immune system, then musical roles may behave according to detection, recognition, escalation, tolerance, memory, and autoimmune failure.

No white-blood-cell lyrics required.

The rules are the import.

Not the imagery.

There is another mechanism where two categorically incompatible things are temporarily declared to be literally the same thing.

A is B.

No metaphor.

Then the system keeps reasoning normally and has to absorb the consequences.

Another one creates multiple incompatible perceptual readings and refuses to average them into consensus.

Instead the weirdest legitimate reading becomes sovereign while the others remain active as resistance constraints.

Another one changes the idea of conceptual distance.

Normally an AI associates things because they are semantically similar.

This system might define similarity according to:

how they fail
what residue remains after destruction
what maintenance burden they impose
how they collapse

what absence they create when removed

Now two concepts that are normally very far apart can become near neighbors under that temporary ruler.

That gives me associations I probably would not have reached through ordinary genre or mood prompting.

There are Lil Guys that selectively remove one cognitive operation and force the remaining reasoning to build a prosthesis around the missing ability.

There are ones that invent synthetic senses.

For example:

detect structural instability

encode it as heat

when heat increases, shorten phrase length

if heat becomes concentrated, reroute melodic responsibility to another instrument

Now the AI is effectively composing through a temporary artificial sensorium.

There are synthetic emotions too.

Not “make it sad.”

Actually define a trigger, detection system, intensity curve, and behavioral consequence.

Maybe the temporary emotion is an intense attraction to unresolved asymmetry.

Whenever a musical structure becomes almost symmetrical but fails, the system becomes increasingly attached to it and protects it from resolution.

That synthetic emotion changes selection.

It is not prose flavor.

There is LOSSY COGNITION.

That deliberately destroys information.

The system compresses two unrelated musical ideas into a crude representation where they become indistinguishable.

Then it reconstructs the music without restoring all of the information that was thrown away.

Artifacts from the collision survive.

There is PROCEDURAL TABOO.

The direct route to the goal is forbidden.

The obvious workaround is also forbidden.

Then the system has to perform some independently valid operation for another reason and allow the desired musical result to emerge only as aftermath.

There is ALIEN UTILITY.

The prompt can define a reward function that values something unusual.

Maybe it rewards structures that create useful isolated stems.

Maybe it rewards transformations that preserve identity while maximizing internal disagreement.

Maybe it rewards future mutability instead of immediate beauty.

Then the system generates several valid possibilities and selects according to that temporary utility function instead of generic “good music” preference.

There is DESCENDANT FITNESS.

Instead of asking which idea is best right now, the system imagines several generations of descendants from each idea and chooses the ancestor that creates the richest future territory.

That one is especially useful because I am not trying to produce one perfect song.

I am trying to create systems that keep producing new musical territory.

There is RECONSTRUCTIVE MEMORY.

Every time an earlier musical idea gets recalled, the context in which it is recalled is allowed to scar it.

So the remembered idea slowly mutates over the course of the piece.

There is ATTENTIONAL COLLAPSE.

Once a property is examined closely, it crystallizes and becomes difficult to change, while unexamined properties remain free to mutate.

There is SEMANTIC RECOIL.

A later musical event can change what an earlier event means.

The earlier event itself does not change.

Its interpretation does.

Then all later reasoning has to use the updated interpretation.

And there is a larger breeding system where the cognitive rules themselves can reproduce.

This part is important:

I do not just stack two cognitive mechanisms one after another.

I can ask:

What NEW single operation becomes possible because these two mechanisms interact?

Then the resulting offspring rule has to contain ancestry from both parents while being reducible to neither.

If parent A alone could have produced it, it fails.

If parent B alone could have produced it, it fails.

If it is just A followed by B, it fails.

The offspring has to be a genuinely new operation.

Then that new cognitive rule gets used to generate music.

And once it generates something, the output changes the conceptual environment.

That changes which rules are useful next.

So the system can go:

rule population

→ hybrid cognitive offspring

- musical idea
- changed conceptual environment
- changed fitness pressure
- new recombination

That is why I say the Lil Guys are integrated.

They are not a menu of random prompt tricks sitting beside each other.

They plug into a common architecture.

The cognitive layer changes HOW candidate musical ideas are generated.

The jurisdiction layer decides WHICH musical dimension each idea gets authority over.

The operator layer decides HOW those systems transform over time.

The anchor layer determines WHAT has to remain legible.

The arrangement layer determines WHEN systems appear, disappear, interrupt each other, or respond to events.

The cast layer determines WHO is allowed to do what.

The feedback layer records which mechanisms actually produced results I liked.

So the system is hierarchical but connected.

A cognitive rule might generate the idea that rhythmic density should behave like a contagious process.

Then the rhythm jurisdiction owns that behavior.

Then an operator might say the contagion can cross into vocals only after a specific threshold.

Then the cast architecture might assign one crowd population to become the carrier.

Then the event-driven arrangement might say that once enough voices are infected, the harmony disappears entirely.

Then the anchor might be a four-note whistle that continues unchanged through the whole catastrophe.

Now all of those pieces are doing different jobs, but they belong to one machine.

That is the part I find really exciting.

The system is not just “more random knobs.”

The knobs interact through structure.

A lot of this also came from analyzing the Suno songs that actually worked for me.

And a very obvious pattern emerged:

I do not really like chaos for its own sake.

I like several highly legible systems becoming excited at each other.

That distinction explains a LOT of the music.

The songs I like tend to have something the nervous system can grab onto:

- a pulse
- a bass pocket
- a repeated phrase
- a narrator
- a tiny melodic hook
- a vocal role
- a recognizable riff
- a stable physical texture

Then the rest of the composition is allowed to become increasingly unhinged around that anchor.

That led to a set of behavior recipes.

Not genres.

Behavior starters.

For example:

JURISDICTION LAB

Harmony, melody, rhythm, timbre, vocals, and performance behavior obey separate laws. Entire jurisdictions may temporarily disappear while the others are forced to continue without replacing them.

PANIC ENGINE

The defining parameter is not just tempo. It is EVENT ARRIVAL RATE.

Something musically distinct happens before the previous thing has completely settled.

Announcement.

Answer.

Instrumental interruption.

Vocal relay.

Meter change.

Diagnostic intrusion.

Mass shout.

Hard stop.

The density is made from identifiable actors rather than undifferentiated mush.

NARRATOR + MOB

Start with one dry, intelligible, emotionally specific voice.

Then introduce secondary vocal populations with completely different jobs.

Speech becomes song.

Song becomes percussion.

Percussion becomes chant.

Chant fractures into hocketed individual voices.

Everything vanishes except one stupid human noise.

CAST OF FREAKS

Instead of “lead singer plus backing vocals,” every vocal population is treated as a separate musical entity.

HOST

LEAD

AUDIENCE

CHOIR

ROBOT

SMALL GROUP

FREAK VOICE

BAND

Each has a jurisdiction.

Eventually they may infect one another's jobs.

EVENT-DRIVEN VARIETY SHOW

The form is not:

verse
chorus
verse
chorus
bridge

The form is:

EVENT HAPPENS
→ MUSIC RESPONDS

Buzzer → harmony freezes.
Reveal → brass enters.
Shout → percussion doubles.
Error → meter changes.
Host interruption → band hard-stops.
Crowd agrees → close harmony appears.
System failure → every stem except one disappears.

That is one of my favorite compositional ideas because it treats arrangement as a responsive system.

ANCHOR UNDER SIEGE

Keep one tiny invariant alive while everything around it undergoes increasingly violent transformations.

DOUBLE-TIME HALLUCINATION

This came from noticing that BPM and perceived kinetic density are not the same thing.

A song can have a 100 BPM skeleton and a 200 BPM bloodstream.

So I am starting to control things like:

BASE PULSE

ACTIVITY PULSE
VOCAL DENSITY
PERCUSSION SUBDIVISION
EVENT ARRIVAL RATE

separately.

That gives me frantic without just typing FAST SONG.

I am also building STEMMINESS into the system.

This means encouraging arrangements where individual parts remain interesting and separable.

At high stemminess:

different instruments occupy distinct rhythmic jobs
different registers remain separated
shared reverb is reduced
parts appear in exposed windows
mutations happen sequentially instead of everyone vomiting at once
performers have identifiable roles
whole jurisdictions can disappear cleanly

That way I can take a track into Suno Studio and replace a stem with something completely idiotic later.

For example:

replace the guitar with a goose having a nervous breakdown.

Which is, scientifically speaking, important.

Another thing I am adding is feedback.

If I star something because I really like it, I do not just want the system to know:

Merry liked this song.

That is not useful enough.

I want to tell it:

I liked the way the deadpan narrator stayed completely stable while the crowd became increasingly rhythmic.

I liked the 3-against-2 pulse relationship.

I liked that the harmony disappeared before the final section.

I liked that the whistle stayed recognizable through all the mutations.

I liked that the vocals became percussion.

I liked that the activity level felt twice as fast as the actual BPM.

Over time, those observations start revealing something like a personal compositional genome.

Not a genre.

Not a "Merry sound."

A set of recurring structural preferences.

So far mine seems to strongly favor:

multiple rhythmic truths

communal vocals

casts rather than backup singers

fast explanatory or patter-style language

hard contrast instead of smooth transitions

small anchors surviving large transformations

dry physical textures

stomps, claps, cheap organs, whistles, brass attacks, fuzz, spring reverb, hand percussion

performers that structurally alter the arrangement

high kinetic density

legible weirdness

and, very importantly:

EXCITEMENT WITHOUT MUSICAL STUPIDITY.

I do not just want “happy upbeat.”

I want the piece to be intellectually doing shit while losing its fucking mind.

That is the actual project.

I am using Suno less like a magic jukebox and more like an unpredictable instrument attached to a weird composition laboratory.

The question is not really:

“What song do I want?”

The question is:

“What rules should govern this temporary musical universe, and what happens when those rules are allowed to interact?”

Sometimes the result is beautiful.

Sometimes it is deeply wrong.

Sometimes it sounds like six bureaucrats, a barbershop quartet, a broken carnival ride, a wildlife narrator, and an accordion arguing about topology inside a laundromat.

But if something weird happens because a specific rule caused it, then even the failures are useful.

Because now I have information.

I can keep the mechanism.

Change the mechanism.

Breed the mechanism with another one.

Move it into a different musical jurisdiction.

Increase or decrease its authority.

Change what triggers it.

Change what survives it.

Or save it as part of the growing library.

So the real thing I am building is not a prompt generator.

It is a system for generating and testing compositional thought.

The songs are the evidence.